

- I. Introduction
 - A. Brief Intro- System call trace recording and replaying
 - B. Project Motivation
 - 1. Accurately record and replay workloads using Linux system calls
 - 2. Uses:
 - a) Benchmarking operating system features/alterations
 - b) Analyzing program security and other properties
- II. Workflow
 - A. Strace -> library -> Dataseries file -> replayer
 - 1. Minimize changes to strace
 - 2. No changes to DS library
 - 3. Glue code (C to C++) in helper "strace2ds" library
 - B. High level overview (diagram).
- III. Recording
 - A. Why our approach?
 - 1. Problems with old approach (csv)
 - B. Approach:
 - 1. Minimize changes to strace code
 - 2. Every change wrapped in `#ifdef ENABLE_DATASERIES`
 - a) Enabled with `./configure --enable-dataseries`
 - 3. Run-time option `"-X <ds-filename>"` needed to turn functionality at run time.
 - 4. When writing to DS file, normal strace output (e.g., pretty-printing) is off
 - 5. Strace2ds library also uses `autoconf`
 - C. Where/how we get information in strace
 - 1. Struct `tcb`
 - a) Most relevant information changed for each system call
 - 2. Common fields/`v_args`
 - a) Void * arrays
 - b) Common fields:
 - (1) Time called, time returned, `retval`, `errno`, `pid`
 - c) `V_args`
 - (1) Virtual arguments: to copy mem from `ptrace'd` process to strace's addr space
 - (2) Used to pass pathnames, read/write buffers, `stat` structures, etc. for system calls that require them
 - 3. Our utility functions in `util.c`
 - a) `ds_get_*`
 - b) Get pathnames/different types of buffers for different system calls
 - c) Call strace's `umoven` function, which copies data from the address space of process being traced to that of strace

4. Switch block in syscall.c
 - a) tcp->s_ent->sen
 - (1) Strace's own identifying number for system calls
 - (2) I.e. SEN_open
 - b) Mostly in trace_syscall_exiting()
 - (1) Called after actual system call is executed
 - (2) Fields in struct tcb contain accurate values for that system call at that point
 - c) Non-terminating system calls in trace_syscall_entering()
 - (1) _exit(), execve()
 - (2) These calls either never reach trace_syscall_exiting() or the values passed as arguments are inaccessible from trace_syscall_exiting()

D. What happens in the library

1. Wrapper code (strace2ds.h, strace2ds.cpp)
2. Fields table
3. XML files (created from generate-xml.sh <tablefile>)
4. DataSeriesOutputModule
 - a) One DSOM object per trace
 - (1) When initialized: creates configuration table with extent names (syscall names) and their relevant fields
 - b) Walk through what happens for each system call
 - (1) ds_write_record(...)
 - (2) Creates map of field names to field values (string to void *)
 - (3) Common fields stored in map (if present)
 - (4) Check if syscall name matches a supported call
 - (5) Make[syscall name]ArgsMap function
 - (a) Stores system call's specific fields/arguments in map
 - (6) Iterate through field names and write to DataSeries file
 - c) Any specific system calls we want to point out?

E. How to run (as an option of strace)

1. STRACE2DS=~/.strace2ds ./strace -X foo.ds <executable>

IV. Replaying

A. Approach

1. Base SystemCallTraceReplayModule class
2. Individual system call module classes
3. Priority Queue

B. Workflow

1. Initialize a module for each supported system call
2. Replay in order of unique_id number
3. processRow(): defined in each derived class
 - a) Gets argument values from the ExtentSeries
 - b) Actually replays system call
4. completeProcessing():
 - a) after_sys_call()
 - (1) Compares retval, errno
 - (2) Prints system call fields if desired
 - b) Adjusts series location
 - (1) (not actually a pointer, but ++operator is overloaded in DataSeries so that we can move to the next row similarly)

C. Replaying Options

1. Default
 - a) Prints a message when the first of one system call (i.e., close) is played, and when the last of that system call is played.
2. Verbose
 - a) Prints each system calls common fields/arguments
3. Verify
 - a) Verifies traced and replayed read/write, stat, getdents buffers contain the same data
4. Warn
 - a) Prints a warning message if recorded and replayed retval/errno aren't the same
5. Abort
 - a) Aborts replayer if recorded and replayed retval/errno aren't the same
6. Write pattern data
 - a) Rand() or dev/urandom
 - b) Repeated pattern
 - c) 0s

D. Replaying Design Decisions

1. File descriptor map
 - a) Maps recorded fd to replayed fd
 - b) Map certain standard values prior to replay
 - (1) STDIN, STDOUT, STDERR, AT_FDCWD
2. Using integer encoding of mode/flag values
 - a) Instead of recovering these values from the boolean flag/mode fields specified in the SNIA doc (faster)
3. Rows per call (readv, writev, execve)

- a) Most system calls require one record and take up one row in an extent
 - (1) Default value of rows_per_call is 1
 - b) Some require more than one
 - c) Rows_per_call will be set accordingly in processRow
 - 4. System calls that don't make sense to replay (but we record them)
 - a) Ex: exit, execve
 - E. How to RUN syscall-replayer
 - 1. ./system-call-replayer <foo.ds>
- V. System calls supported
 - A. list
- VI. Programs/utilities we have traced and replayed successfully
 - A. Cp, mv, rm, ls, ...
- VII. Changes to SNIA doc
 - A. All changes redlined in DOC file
 - B. Added some system calls
 - 1. Rename, getdents, openat, unlinkat
 - 2. Added/removed/ certain fields
 - 3. Made certain fields (non-)nullable
 - 4. Fixed typos and inconsistencies
- VIII. Other issues not yet covered
 - A. Testing
 - 1. Short test programs for individual system calls
 - B. Stats we need (e.g., timing replayed syscalls)
 - C. More syscalls we know we need to capture (mmap group, exec*, clone, ioctl)
- IX. To do next:
 - A. Strace will complain if it traces a syscall we're not capturing
 - B. Replay bigger and bigger apps
 - 1. Ultimate goal: replay server apps (mysql, apache, etc.)